

**NANYANG**  
**TECHNOLOGICAL**  
**UNIVERSITY**

**CZ4046 Intelligent Agents**  
**Assignment 1 – Agent Decision Making**

Jervis Chan Jun Yong  
U1821352H

## Table of Contents

<b>1. INTRODUCTION .....</b>	<b>3</b>
<b>2. VALUE ITERATION.....</b>	<b>5</b>
<b>DESCRIPTION OF IMPLEMENTATION .....</b>	<b>5</b>
<b>RESULTS.....</b>	<b>7</b>
<i>Plot of Optimal Policy.....</i>	<i>7</i>
<i>Utilities of All States.....</i>	<i>8</i>
<i>Plot of Utility Estimates as a Function of Number of Iterations.....</i>	<i>8</i>
<i>Studying Effects of <math>\epsilon</math> Against Convergence Iteration.....</i>	<i>9</i>
<b>3. POLICY ITERATION .....</b>	<b>10</b>
<b>DESCRIPTION OF IMPLEMENTATION .....</b>	<b>10</b>
<b>RESULTS.....</b>	<b>12</b>
<i>Plot of Optimal Policy.....</i>	<i>12</i>
<i>Utilities of All States.....</i>	<i>13</i>
<i>Plot of Utility Estimates as a Function of Number of Iterations.....</i>	<i>14</i>
<b>4. BONUS QUESTIONS .....</b>	<b>15</b>
<b>DESCRIPTION OF IMPLEMENTATION .....</b>	<b>15</b>
<b>QUESTION: HOW DOES THE NUMBER OF STATES AFFECT CONVERGENCE .....</b>	<b>16</b>
<b>QUESTION: HOW DOES INCREASING COMPLEXITY AFFECT CONVERGENCE .....</b>	<b>18</b>
<b>5. SOURCE CODE .....</b>	<b>19</b>
<b>6. CONCLUSION .....</b>	<b>20</b>

## 1. Introduction

In this assignment, we are tasked to solve a Markov Decision Process (MDP). The MDP is composed as follows:

### States - s

We are provided with a 6 x 6 grid environment (Figure 1.1).



Figure 1.1. 6 x 6 Grid Environment

There are four kinds of squares in the grid: green, white, orange and wall. As an agent cannot be in the wall, there are 31 possible states. The states are defined by coordinates which are in the (row, col) format, with the top left corner being (0,0).

### Actions - A(s)

In this maze, there are four possible actions, namely up, down, left and right. At each state, an agent can take three of the four possible actions. It can move in its intended direction or at right angle to the intended direction. If the move makes the agent walk into a wall, the agent will stay in the same place.

### Transition model - $P(s' | s, a)$

The transition model is stated as follows (Figure 1.2). The transition properties depend only on the current state and not on the previous history, demonstrating the Markov Property.




Figure 1.2. Transition Model of Agent

### Reward function - R(s)

The reward function gives us the reward given when the agent is a certain state.

$$R(\text{white square}) = -0.04$$

$$R(\text{green square}) = +1$$

$$R(\text{brown square}) = -1$$

### Policy $\pi(s)$

The policy  $\pi(s)$  gives us the action that the agent will take in any given state. An optimal policy should maximise the expected utility over all possible state sequences. In this assignment, we aim to find the optimal policy to the MDP using **value iteration** and **policy iteration**.

As there are no terminal states, the agent's state sequence could be infinite. Thus, we will use a discount factor ( $\gamma$ ) of 0.99 to ensure that the total utility will stay bounded and allow the algorithms to converge.

## 2. Value Iteration

### Description of Implementation

In value iteration, we calculate the utility of each state and use the state utilities to decide on the optimal action in each state.

The implementation can be described in the following steps:

1. Create dictionaries which will store current and next utility values of each state in the MDP. The utility values of each state are initialised to 0.
2. Initialise variable which will hold the maximum change in utility at any iteration.
3. Calculate convergence threshold using the following formula:

$$\text{Convergence Threshold} = \frac{\epsilon (1 - \gamma)}{\gamma}$$

$\epsilon$  represents Epsilon, which is the maximum allowed error in utility. It can be set at different values to see how the convergence of value iteration will be affected.

4. Iterate till convergence
  - a. Perform Bellman Update for each state to calculate utility for the next state and set action that maximises expected utility. The equation for the Bellman Update is as follows:
$$U_{i+1}(s) \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U_i(s')$$
  - b. Calculate maximum change in utility among all states and compare against convergence threshold. Convergence is reached when maximum change is less than the convergence threshold.

### Code Implementation of Value Iteration

The above implementation can be seen in the code snippets below. Three functions are created:

1. *value\_iteration* (Figure 2.1)  
This function is the main code implementation of the value iteration algorithm above.
2. *bellman\_update* (Figure 2.2)  
This function is the implementation of the Bellman Update formula. For each action, the function will call *get\_expected\_utility* to calculate the utility of that action and determine what is the action that will return the maximum utility.
3. *get\_expected\_utility* (Figure 2.3)  
This function will be given the current state and action and returns the expected utility.

```

1. def value_iteration(mdp, EPSILON):
2.     # LOCAL VARIABLES
3.     U_current, U_next = {}, {} # U, U' - dictionary of utilities for states
4.     max_utility_change = 0 #  $\delta$  - maximum change in utility at any iteration
5.     convergence_threshold = (EPSILON * (1 - mdp.discount))/mdp.discount
6.
7.     # Supplementary dictionaries for visualisation
8.     U_iterations, optimal_policy, policy_loss= {}, {}, []
9.
10.    for state in mdp.states:
11.        #initialise utility value of all states in U_current and U_next to be 0
12.        U_current[state] = 0
13.        U_next[state] = 0
14.        U_iterations[state] = []
15.        optimal_policy[state] = None
16.
17.    num_iterations = 0
18.    converged = False
19.
20.    while not converged:
21.        #  $U \leftarrow U'$ ;  $\delta \leftarrow 0$ 
22.        for state in mdp.states:
23.            U_iterations[state].append(U_current[state])
24.            U_current[state] = U_next[state] #  $U \leftarrow U'$ 
25.
26.        max_utility_change = 0 #  $\delta$ 
27.
28.        for state in mdp.states:
29.            updated_utility, best_action = bellman_update(mdp, state, U_current)
30.            U_next[state] = updated_utility
31.            optimal_policy[state] = best_action
32.
33.            abs_utility_change = abs(U_next[state] - U_current[state])
34.            if abs_utility_change > max_utility_change:
35.                max_utility_change = abs_utility_change
36.
37.        num_iterations += 1
38.        policy_loss.append(max_utility_change)
39.        if max_utility_change < convergence_threshold:
40.            converged = True
41.
42.    return {'U_current': U_current, 'U_iterations': U_iterations,
'optimal_policy': optimal_policy, 'num_iterations' : num_iterations, 'policy_loss':
policy_loss}

```

Figure 2.1. Code Snippet of value\_iteration Function

```

1. def bellman_update(mdp, state, U_current):
2.     """
3.     Bellman Update:  $U_{i+1}(s) \leftarrow R(s) + \gamma \max_a \sum P(s'|s,a) U_i(s')$ 
4.
5.     """
6.     max_expected_utility = float('-inf')
7.     best_action = None
8.
9.     for action in mdp.actions:
10.        expected_utility = get_expected_utility(mdp, state, action, U_current)
11.        if expected_utility > max_expected_utility:
12.            max_expected_utility = expected_utility
13.            best_action = action
14.
15.    utility = mdp.reward_function(state) + mdp.discount * max_expected_utility
16.    return (utility, best_action)

```

Figure 2.2. Code Snippet of bellman\_update Function

```

1. def get_expected_utility(mdp, state, action, U_current):
2.     """
3.     Calculates sum of expected utility:  $\sum s'P(s'|s, a)U[s']$ 
4.     """
5.     sum_expected_utility = 0
6.     next_states = mdp.get_next_states(state, action)
7.
8.     for next_state in next_states:
9.         probability = mdp.transition_model(state, action, next_state)
10.        actual_next_state = next_states[next_state]["actual"]
11.        sum_expected_utility += probability * U_current[actual_next_state]
12.
13.    return sum_expected_utility

```

Figure 2.3. Code Snippet of get\_expected\_utility Function

## Results

The following plots are obtained using  $\epsilon$  value of 1 and  $\gamma$  of 0.99. With these values, convergence was reached after 459 iterations.

### Plot of Optimal Policy

Figure 2.4 shows the final plot of the optimal policies of all states after convergence of the value iteration algorithm.




Figure 2.4. Value Iteration Optimal Policy Plot with  $\epsilon = 1$

As seen from Figure 2.4, the optimal policy given the starting point of (3,2) is to move left then move to the top towards the reward.

### Utilities of All States

Figure 2.5 shows the final policies of all states after convergence of the value iteration algorithm.

Value Iteration Algorithm		(2, 1) : 94.58430914539592
Values of Parameters		(2, 2) : 92.2923090083678
		(2, 3) : 92.17415441594787
		(2, 4) : 92.10025046569073
		(2, 5) : 90.79275246663413
Discount Factor	: 0.99	(3, 0) : 94.5517204953327
Epsilon Value	: 1	(3, 1) : 93.45037519604443
Number of Iterations	: 459	(3, 2) : 92.23042681579794
Utility Values of all States		(3, 3) : 90.11313792435135
		(3, 4) : 90.8122884668515
		(3, 5) : 90.88596595530619
State : Value		(4, 0) : 93.31040080578784
(0, 0) :	98.9978813938348	(4, 4) : 88.54629449401432
(0, 2) :	94.04333862798383	(4, 5) : 89.56464706487044
(0, 3) :	92.872882390745	(5, 0) : 91.93535571082931
(0, 4) :	91.65249584464884	(5, 1) : 90.72665902358499
(0, 5) :	92.32638442714297	(5, 2) : 89.53303336735367
(1, 0) :	97.3912429045024	(5, 3) : 88.35429082404296
(1, 1) :	94.8808987787996	(5, 4) : 87.56698047101396
(1, 2) :	93.54287976274642	(5, 5) : 88.29557198218212
(1, 3) :	93.39559622989613	
(1, 5) :	89.91580458882243	
(2, 0) :	95.94638157569194	

Figure 2.5. Final Utility of All States with  $\epsilon = 1$

### Plot of Utility Estimates as a Function of Number of Iterations

Figure 2.6 shows the plot of utility estimates against number of iterations of the value iteration algorithm.

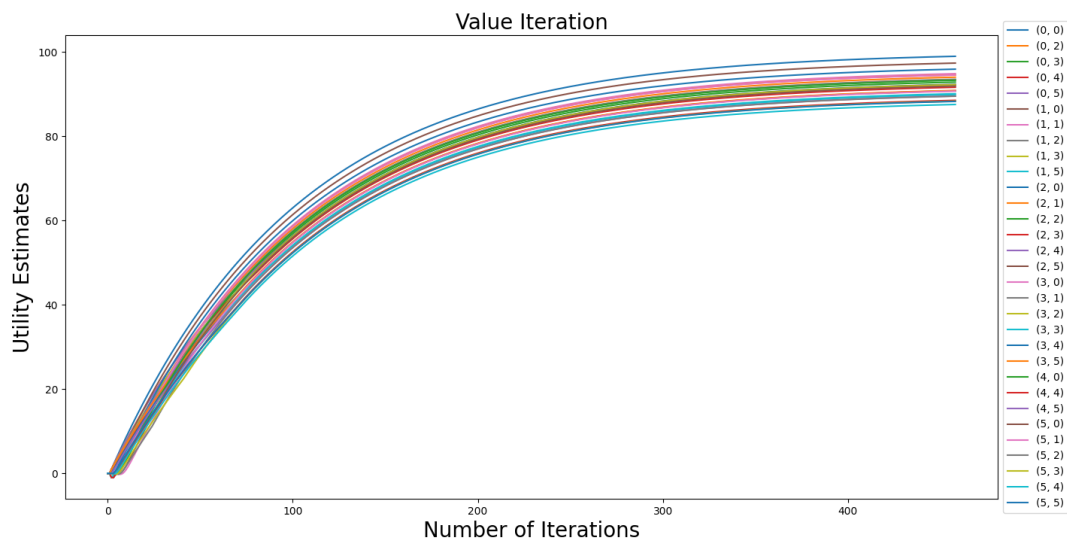


Figure 2.6. Plot of Utility Estimates Against Number of Iterations



### Studying Effects of $\epsilon$ Against Convergence Iteration

In addition to performing value iteration to obtain the optimal policy, we studied the effects of  $\epsilon$  on convergence. This was done by varying the value of  $\epsilon$  was varied and seeing how many iterations it takes to achieve convergence.




Figure 2.7. Plot of Utility Estimates Against Number of Iterations

As seen from Figure 2.7, reducing  $\epsilon$  will help to reduce the number of iterations it takes for the value iteration to converge. This directly relates to allowing a larger margin of error. However, this larger margin of error could also result in incorrect policies. Thus, more experimentation could be done within that range to find the optimal  $\epsilon$  which can balance the correctness of the results and the time taken for convergence.




Figure 2.8. Plot of Optimal Policy,  $\epsilon = 15, 30, 60$

As seen from Figure 2.8, the optimal policies for all states for  $\epsilon = 15$  and 30 are the same as  $\epsilon = 1$ . However, when  $\epsilon$  gets too large, such as  $\epsilon = 60$ , we can observe differences in the optimal policy among some states.

### **3. Policy Iteration**

#### **Description of Implementation**

In policy iteration, we start from an initial policy  $\pi_0$ , and alternate between **Policy Evaluation** and **Policy Improvement** until Policy Improvement yields no change in utilities.

In **policy evaluation**, we are given a policy  $\pi_i$ , and we calculate the utility of each state if that policy were to be executed. As the given policy helps specify the action in the current state, we can use a simplified Bellman Update as follows:

$$U_{i+1}(s) \leftarrow R(s) + \gamma \sum_{s'} P(s' | s, \pi_i(s)) U_i(s')$$

This is repeated  $k$  times to produce the next utility estimate. The value of  $k$  can be experimented to see how it affects convergence.

In **policy improvement**, we take the utilities calculated in policy evaluation and determine whether to change the policy by comparing if the max utility of all possible actions is higher than the utility calculated with the current policy.

The implementation can be described in the following steps:

1. Create dictionaries which will store utility values and policies of each state in the MDP. The utility values of each state are initialised to 0 and the initial policy is randomly selected.
2. Initialise *unchanged* Boolean variable to False to flag changes to policy.
3. Iterate till convergence
  - a. Perform policy evaluation which calculates utility for each state using the current policy.
  - b. Pass the utility values into policy improvement which will return new policy if there is a better policy among all possible actions. Convergence is reached when there is no change from policy improvement.

#### **Code Implementation of Policy Iteration**

The above implementation can be seen in the code snippets below. Three functions are created:

1. *policy\_iteration* (Figure 3.1)  
This function is the main code implementation of the policy iteration algorithm above.
2. *policy\_evaluation* (Figure 3.2)  
This function is the implementation of policy evaluation. The function will perform the simplified Bellman Update  $k$  times using a loop and return the updated policy values.
3. *policy\_improvement* (Figure 3.3)  
This function is the implementation of policy improvement. The function will iterate through the different possible actions and determine the best possible action. It will then compare the utility value of the best action and the current policy and determine if a policy change has taken place.

```

1.     def policy_iteration(mdp, k):
2.         # LOCAL VARIABLES
3.         # U , a vector of utilities for states in S , initially zero and  $\pi$ , a policy
vector indexed by state, initially random
4.         utilities, policies = {}, {}
5.
6.         # Supplementary dictionaries for visualisation
7.         U_iterations= {}
8.
9.         # Initialise all states to 0 and policies to a random action
10.        for state in mdp.states:
11.            utilities[state] = 0
12.            policies[state] = mdp.actions[random.randint(0,3)] #the number will
correspond to an action
13.            U_iterations[state] = [0]
14.
15.        unchanged = False
16.        num_iterations = 0
17.
18.        while not unchanged:
19.            utilities, new_iteration_utilities = policy_evaluation(policies,
utilities, mdp, k)
20.            policies, unchanged = policy_improvement(mdp, utilities, policies,
unchanged)
21.            num_iterations += 1
22.
23.            for state in mdp.states:
24.                U_iterations[state].extend(new_iteration_utilities[state])
25.
26.        return {'optimal_policy': policies, 'U_current': utilities, 'U_iterations':
U_iterations, 'num_iterations': num_iterations }

```

Figure 3.1. Code Snippet of policy\_iteration Function

```

1.     def policy_evaluation(policies, utilities, mdp, k):
2.         """
3.         Simplified Bellman Update
4.         """
5.
6.         U_updated = {}
7.         new_iteration_utilities = {}
8.
9.         for state in mdp.states:
10.            new_iteration_utilities[state] = []
11.
12.        for i in range(int(k)):
13.            for state in mdp.states:
14.                # Calculate  $\sum s'P(s'|s, a)U[s']$  where a is the action under  $\pi$ 
15.                expected_utility = get_expected_utility(mdp, state, policies[state],
utilities)
16.
17.                #  $U_{i+1}(s) \leftarrow R(s) + \gamma \sum s'P(s'|s, \pi_i(s)) U_i(s')$ 
18.                U_updated[state] = mdp.reward_function(state) + mdp.discount *
expected_utility
19.
20.            for state in mdp.states:
21.                utilities[state] = U_updated[state]
22.
23.        for state in mdp.states:
24.            new_iteration_utilities[state].append(U_updated[state])
25.
26.        return (utilities, new_iteration_utilities)

```

Figure 3.2. Code Snippet of policy\_evaluation Function

```

1. def policy_improvement(mdp, utilities, policies, unchanged):
2.     """
3.     Decide change if  $\max \sum_{a \in A(s)} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ 
4.     """
5.     #dictionary to store updated policies if policy change
6.     updated_policies = policies
7.     unchanged = True
8.
9.     for state in mdp.states:
10.        # 1. Find  $\max \sum_{a \in A(s)} P(s' | s, a) U[s']$ 
11.        max_expected_utility = float('-inf')
12.        best_action = None
13.
14.        for action in mdp.actions:
15.            expected_utility = get_expected_utility(mdp, state, action,
16.            utilities)
17.            if expected_utility > max_expected_utility:
18.                max_expected_utility = expected_utility
19.                best_action = action
20.
21.        # 2. Find  $\sum_{s'} P(s' | s, \pi[s]) U[s']$ 
22.        policy_utility = get_expected_utility(mdp, state, policies[state],
23.        utilities)
24.
25.        # 3. Compare
26.        if max_expected_utility > policy_utility:
27.            updated_policies[state] = best_action
28.            unchanged = False
29.
30.    return (updated_policies, unchanged)

```

Figure 3.3. Code Snippet of policy\_improvement Function

## Results

With policy iteration, the number of iterations to convergence can take a range of values as the initial policy was randomised. The following plots are obtained using  $k$  value of 5, which reached convergence was after 10 iterations.

### Plot of Optimal Policy

Figure 3.3 shows the plot of utility estimates against number of iterations of the value iteration algorithm.




Figure 3.3 Policy Iteration Optimal Policy Plot with  $k = 5$

As seen from Figure 3.3, the optimal policy for policy iteration, given the starting point of (3,2), is to move left then move to the top, towards the reward. This provides the same policy as value iteration.

### Utilities of All States

Figure 3.4 shows the plot of utility estimates against number of iterations of the policy iteration algorithm. Depending on the number of iterations, the final utility value might vary even with the same  $k$  value. Thus, in policy iteration, it is more important to look at optimal policy.

-----		(2, 1) : 32.56076260485173
Policy Iteration Algorithm		(2, 2) : 30.773819783712888
Values of Parameters		(2, 3) : 31.707255855861053
-----		(2, 4) : 31.980910592356647
Discount Factor	: 0.99	(2, 5) : 30.870374249477912
K	: 5	(3, 0) : 32.52817394924986
Number of Iterations	: 10	(3, 1) : 31.42682860665801
-----		(3, 2) : 30.270681601096282
Utility Values of all States		(3, 3) : 29.598894276044817
-----		(3, 4) : 30.767802549655794
State : Value		(3, 5) : 31.22152306681961
(0, 0) :	36.97433489268308	(4, 0) : 31.28685420249932
(0, 2) :	33.9374417360702	(4, 4) : 28.628460573881963
(0, 3) :	32.798051747981255	(4, 5) : 29.959813481175825
(0, 4) :	31.83189378084475	(5, 0) : 29.911808848752177
(0, 5) :	32.986166810811454	(5, 1) : 28.70311170544192
(1, 0) :	35.36769640017682	(5, 2) : 27.509484972231025
(1, 1) :	32.85735226331835	(5, 3) : 26.75029313629987
(1, 2) :	32.79568847117196	(5, 4) : 27.79993148999396
(1, 3) :	32.93451704466156	(5, 5) : 28.76174576319313
(1, 5) :	30.63488629954799	-----
(2, 0) :	33.92283506021069	

Figure 3.4. Final Utility of All States with  $k = 5$

### Plot of Utility Estimates as a Function of Number of Iterations

Figure 3.5 shows the plot of utility estimates against number of iterations of the value iteration algorithm.

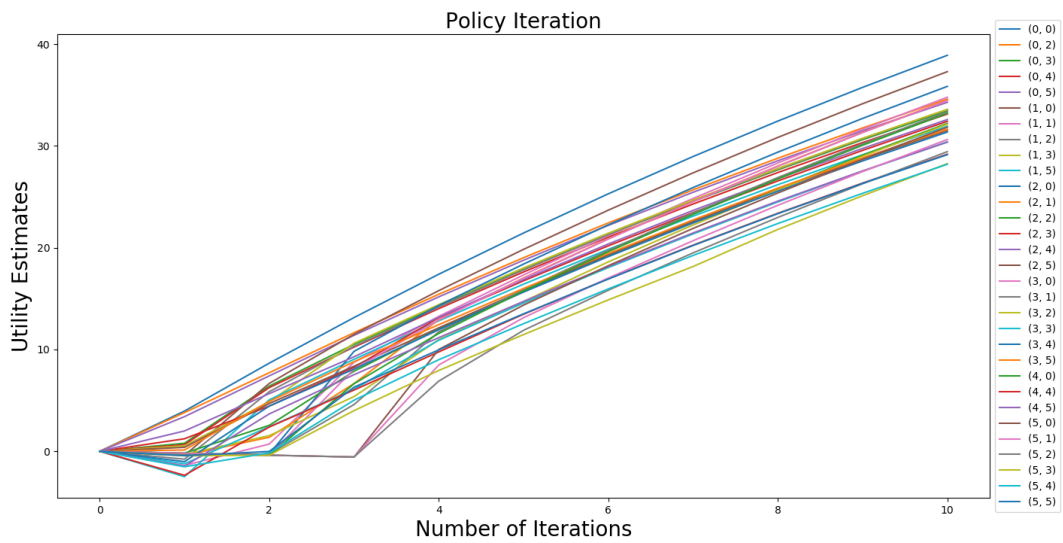


Figure 3.5. Plot of Utility Estimates Against Number of Iterations for  $k = 5$

From Figure 3.5, we can see that the utility estimates against number of iterations for policy iteration follow a different shape compared to utility iteration. In fact, there is even cases where utility values go negative, whilst utility values for policy iteration are strictly increasing from 0. This could be attributed to the fact that random actions were used at the start for policy iteration while value iteration tries to find the best action at every iteration. In fact, as we increase the value  $k$ , we might observe that the shape increasingly becoming like that of the value iteration, with the exception of the first iteration (Figure 3.6).

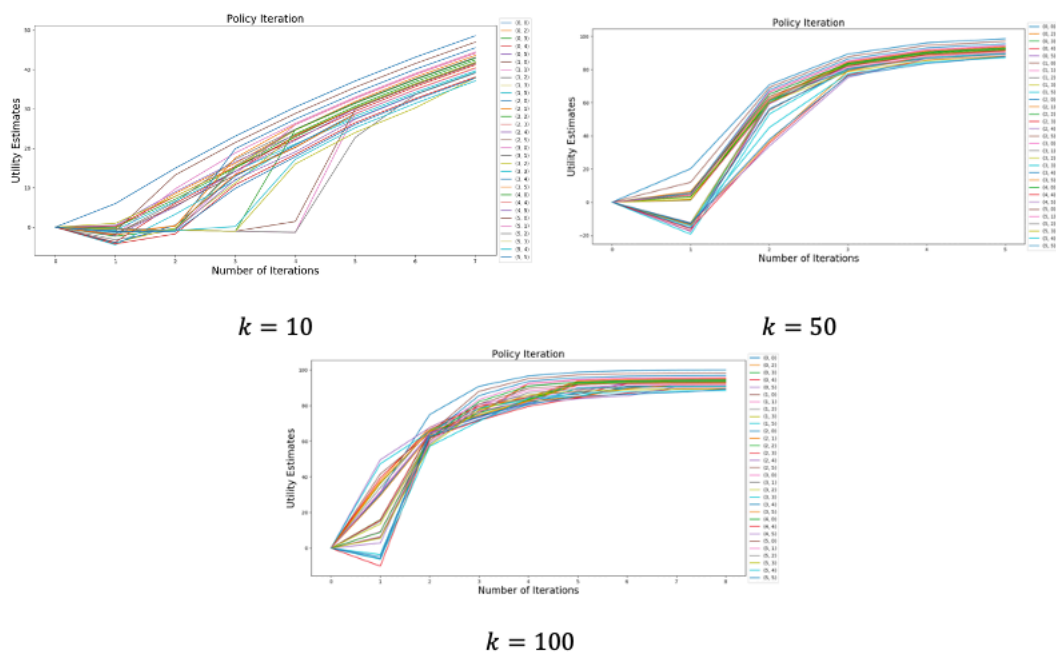


Figure 3.6. Plot of Utility Estimates Against Number of Iterations for  $k = 10, 50, 100$

## 4. Bonus Questions

### Description of Implementation

For the bonus question, the `generate_maze` function (Figure 4.1) was used to generate a more complex maze environment. A more complex maze can be designed by increasing the grid length of the grid or by changing the probability of different types of squares.

This function will be used to generate a maze which will be passed to the value and policy iteration functions. For value iteration,  $\epsilon$  value of 1 was used, while for policy iteration,  $k$  value of 5 was used. These values were picked as they were able to obtain the correct optimal policies.

```
1. def generate_maze(grid_length, special = None, wall = None):
2.     """
3.     Generates a maze with given grid length
4.     """
5.     if special and wall:
6.         w = 1 - special - wall
7.         g = o = special/2
8.
9.     elif special and not wall:
10.        wall = (0.2/0.7) * (1-special)
11.        g = o = special/2
12.        w = (0.5/0.7) * (1-special)
13.
14.    elif wall and not special:
15.        g = (0.15/0.8) * (1-wall)
16.        o = g
17.        w = (0.5/0.8) * (1-wall)
18.
19.    elif not special and not wall:
20.        g = o = 0.15
21.        wall = 0.2
22.        w = 0.5
23.
24.    random.seed()
25.    maze = []
26.    for row in range(grid_length):
27.        maze.append([])
28.
29.        # WE ITERATE THROUGH EACH CELL AND ADD A CELL VALUE
30.        for _ in range(grid_length):
31.            random_color = random.random()
32.
33.            if random_color < w:
34.                maze[row].append(' ')
35.            elif random_color < w + g:
36.                maze[row].append('G')
37.            elif random_color < w + g + o:
38.                maze[row].append('O')
39.            else:
40.                maze[row].append('W')
41.
42.    return maze
```

Figure 4.1. Code Snippet of `generate_maze` Function

### Question: How Does The Number of States Affect Convergence

The number of states available to the agent can be varied using the following methods:

1. Varying number of walls
2. Varying grid length

#### Varying Number of Walls

We change the wall probability provided to *generate\_maze* function to vary the number of walls available in the grid. With more walls, the agent will have less available states to move in. This is implemented by varying the probability of walls while keeping grid size constant. The results from varying wall probability can be seen in Figure 5.1 and 5.2

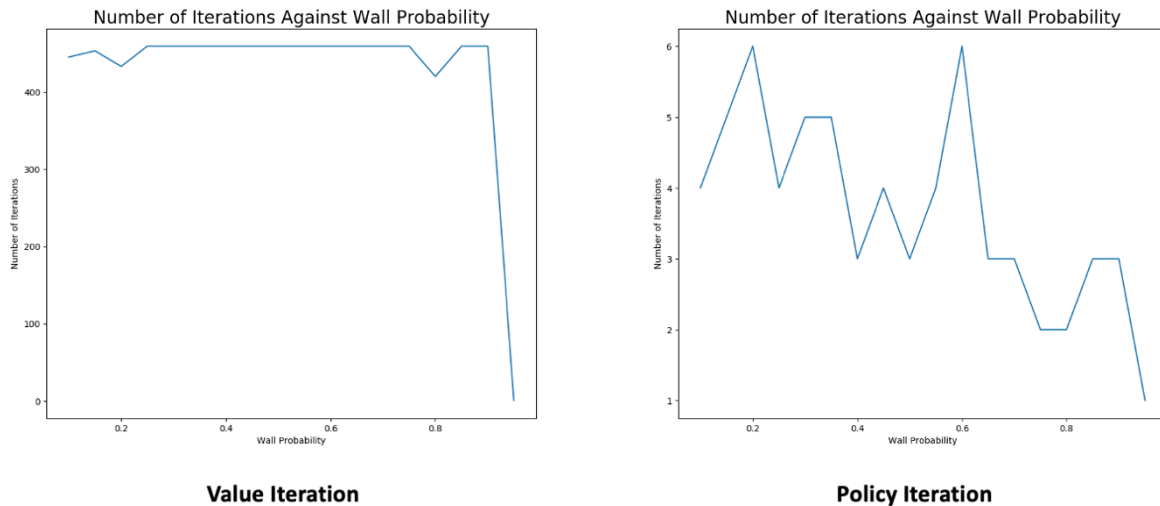


Figure 5.1. Convergence Iterations Against Wall Probability

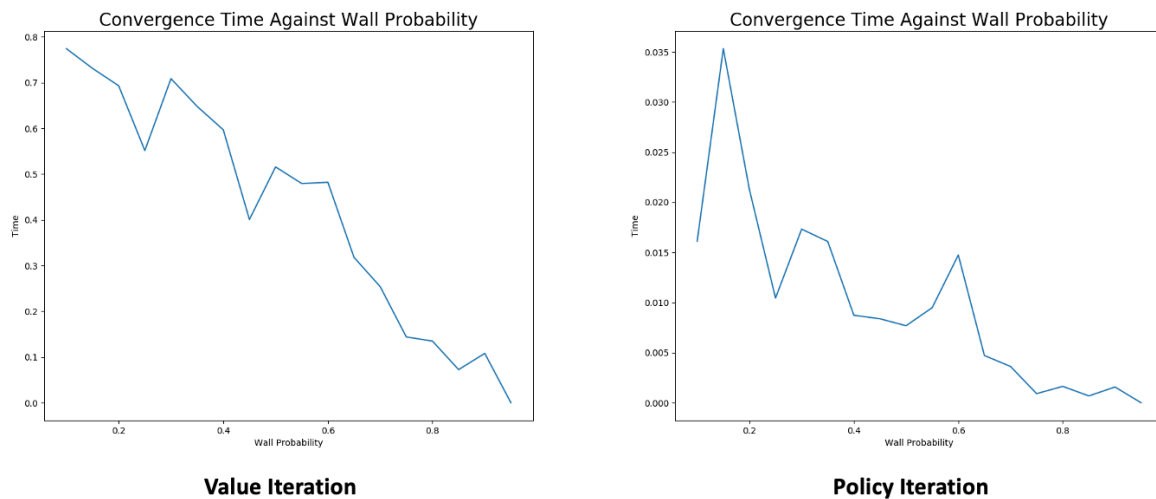


Figure 5.2. Convergence Time Against Wall Probability

In terms of convergence time, we can see that there is a decrease in convergence time as the wall probability (and in turn, number of states) increases. On the other hand, the number of iterations to reach convergence seems to remain the same for value iteration while there seems to be an increasing trend for policy iteration.



### Varying Grid Length

We change the grid length provided to *generate\_maze* function to vary the size of the maze while keeping the probabilities of the different square types forming the same. The results from varying wall probability can be seen in Figure 5.3 and 5.4.

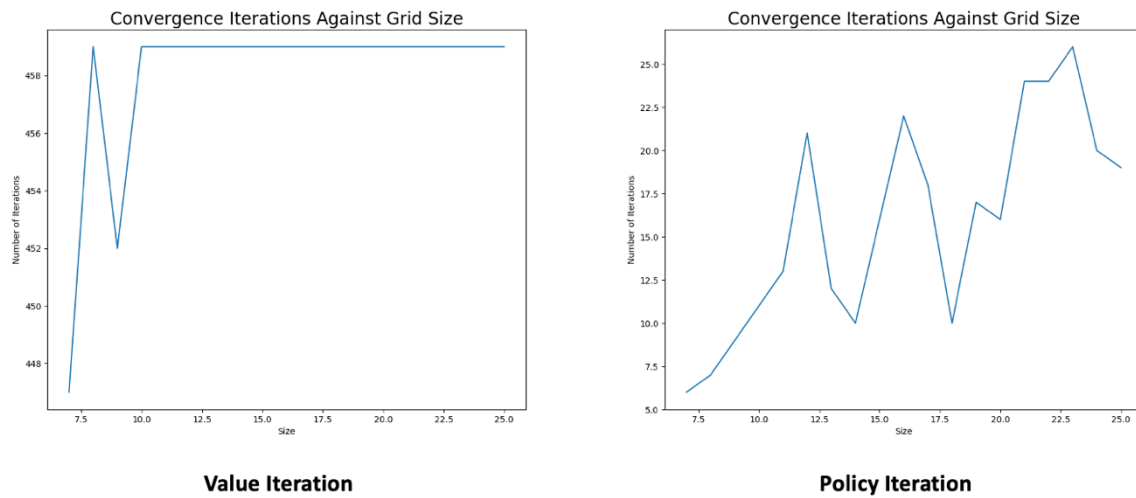


Figure 5.3. Convergence Iterations Against Grid Size

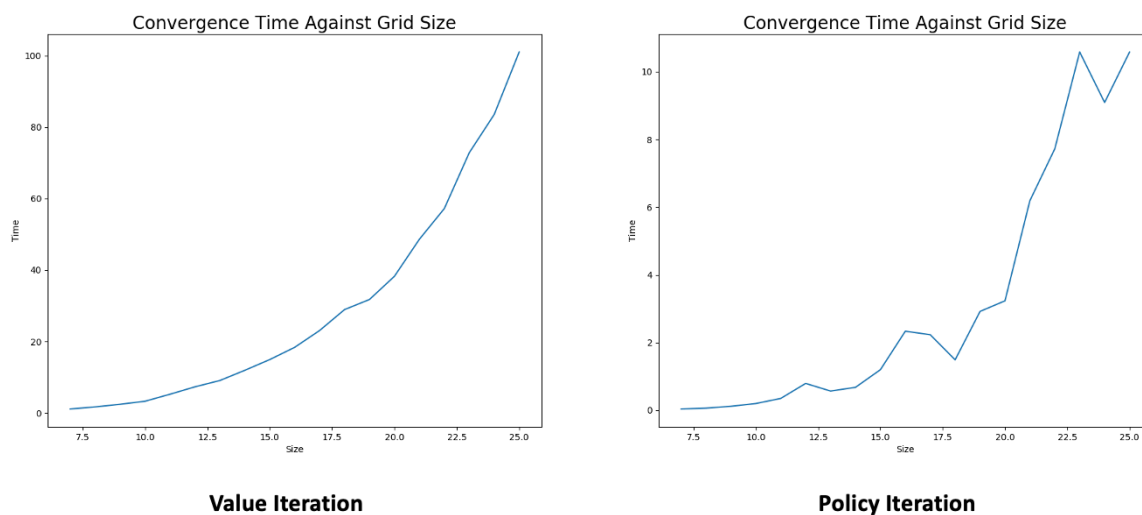


Figure 5.4. Convergence Time Against Grid Size

From Figure 5.3, we can see the number of iterations to reach convergence seems to remain the same for value iteration while there seems to be an increasing trend for policy iteration. On the other hand, in terms of convergence time in Figure 5.4, we can see that there is an increase in convergence time as the grid size (and in turn, number of states) increases. This is the same for both policy and value iteration. However, in terms of absolute time, value iteration takes a much longer time as compared to policy iteration. For instance, with a grid size of 25, the time taken for value iteration is approximately 100s while the time taken for policy iteration is approximately 10s.

For both experiments with number of states, it is more important to look at the effects with regards to convergence time as the trend is clearer. In both experiments, we can see that there is a positive correlation between convergence time and number of states. Additionally, the experiments also show that policy iteration seems better able to scale to handle larger number of states, as compared to value iteration.

### **Question: How Does Increasing Complexity Affect Convergence**

In addition to changing the number of states, we can increase the complexity by varying the number of special squares (rewards and penalties) in the system. This is done by varying the probabilities of the special squares while keeping total grid size the same. Among the special squares, we keep the probability of reward and penalty squares equal. The results from varying wall probability can be seen in Figure 5.5 and 5.6.

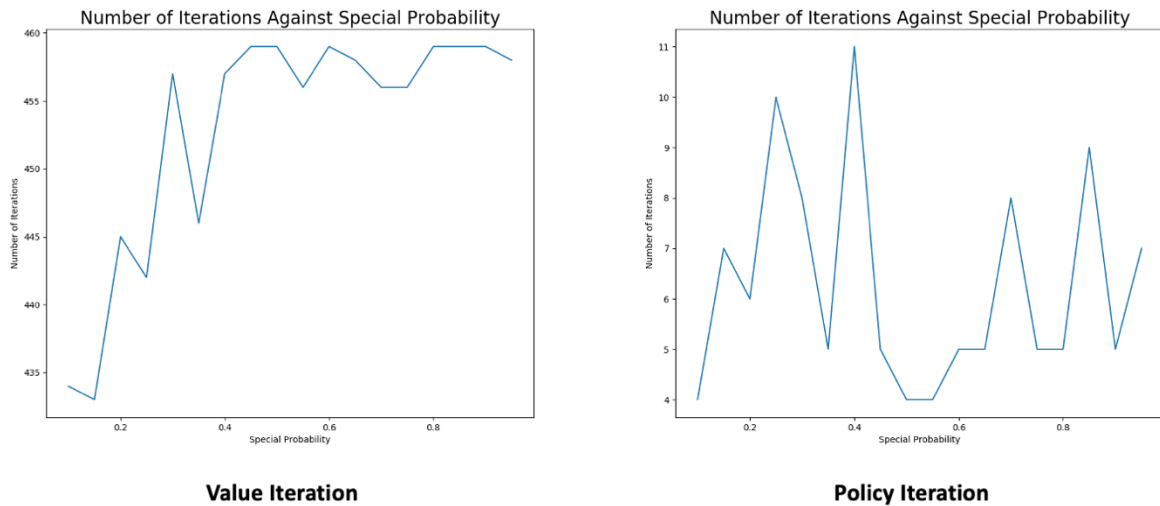


Figure 5.5. Convergence Iterations Against Special Probability




Figure 5.6. Convergence Time Against Special Probability

From Figure 5.5, we can see that there seems to be a positive correlation with number convergence iterations and the probability of special squares for value iteration. On the other hand, from Figure 5.6, there does not seem to be a clear relationship between convergence and probability of special squares. For both algorithms, convergence was reached even as the probability of special squares approach 1. Thus, if we take special squares as an indicator of environment complexity, we can conclude that the algorithms are still able to learn the right policy even an extremely complex environment.

## **5. Source Code**

The source code comprises of the following python files:

1. *main.py*  
The file servers as the manager that will handle the logic required.
2. *constants.py*  
This file holds the constant values which will be used by other files and functions.  
This includes the maze grid, the discount factor and reward map.
3. *MDP.py*  
This file holds the class definition of the abstract MDP class.
4. *Maze.py*  
This file holds the class definition of the Maze class. The Maze will inherit from the MDP class and it is coded according to the MDP defined in the introduction.
5. *valueiteration.py*  
This file holds the implementation of the value iteration algorithm as described in Section 2.
6. *policyiteration.py*  
This file holds the implementation of the policy iteration algorithm as described in Section 3.
7. *bonus\_functions.py*  
This file holds the implementation of the bonus functions to solve the bonus questions as described in Section 4.
8. *visualisations.py*  
This file holds the helper functions to produce the visualisations required.

## **6. Conclusion**

In this assignment, we implemented both value iteration and policy iteration to find the optimal policy in the given MDP environment. As we use these algorithms in the bonus questions, we can see the effects of varying number of states and complexity on the convergence of these algorithms. One insight we can take away is that both algorithms are effective in converging to a policy even in complex environments. Another insight is that the policy iteration algorithm seems better able to scale with increasing state size of the MDP.